
DISTRIBUTED SYSTEMS (COMP9243)

Lecture 7: Fault Tolerance

Slide 1

- ① Failure
 - ② Reliable Communication
 - ③ Process Resilience
 - ④ Recovery
-

DEPENDABILITY

Availability: system is ready to be used immediately

Reliability: system can run continuously without failure

Slide 2

Safety: when a system (temporarily) fails to operate correctly, nothing catastrophic happens

Maintainability: how easily a failed system can be repaired

Building a dependable system comes down to controlling failure and faults.

FAILURE

Terminology:

Failure: a system fails when it fails to meet its promises or cannot provide its services in the specified manner

Error: part of the system state that leads to failure (i.e., it differs from its intended value)

Slide 3

Fault: the cause of an error (results from design errors, manufacturing faults, deterioration, or external disturbance)

Recursive:

- Failure can be a fault
 - Manufacturing fault leads to disk failure
 - Disk failure is a fault that leads to database failure
 - Database failure is a fault that leads to email service failure
-

TOTAL VS PARTIAL FAILURE

Total Failure:

All components in a system fail

- Typical in nondistributed system

Slide 4

Partial Failure:

One or more (but not all) components in a distributed system fail

- Some components affected
 - Other components completely unaffected
 - Considered as *fault* for the whole system
-

CATEGORISING FAULTS AND FAILURES

Types of Faults:

Transient Fault: occurs once then disappear

Intermittent Fault: occurs, vanishes, reoccurs, vanishes, etc.

Slide 5 **Permanent Fault:** persists until faulty component is replaced

Types of Failures:

Process Failure: process proceeds incorrectly or not at all

Storage Failure: "stable" secondary storage is inaccessible

Communication Failure: communication link or node failure

FAILURE MODELS

Crash Failure: a server halts, but works correctly until it halts

Fail-Stop: server will stop in a way that clients can tell that it has halted.

Fail-Resume server will stop, then resume execution at a later time.

Slide 6 **Fail-Silent:** clients do not know server has halted

Omission Failure: a server fails to respond to incoming requests

- *Receive Omission:* fails to receive incoming messages
- *Send Omission:* fails to send messages

Timing Failure: a server's response lies outside the specified time interval

Response Failure: a server's response is incorrect

- *Value Failure:* the value of the response is wrong
- *State Transition Failure:* the server deviates from the correct flow of control

Slide 7

Arbitrary Failure: a server may produce arbitrary response at arbitrary times (aka *Byzantine failure*)

FAULT TOLERANCE

Fault Tolerance:

→ System can provide its services even in the presence of faults

Goal:

→ Automatically recover from partial failure

→ Without seriously affecting overall performance

Slide 8

Techniques:

→ Prevention: prevent or reduce occurrence of faults

→ Prediction: predict the faults that can occur and deal with them

→ Masking: hide the occurrence of the fault

→ Recovery: restore an erroneous state to an error-free state

FAILURE PREVENTION

Make sure faults don't happen:

- Quality hardware
- Hardened hardware
- Quality software

Slide 9



FAILURE PREDICTION

Deal with expected faults:

- Test for error conditions
- Error handling code
- Error correcting codes
 - checksums
 - erasure codes

Slide 10



DETECTING FAILURE

Failure Detector:

- Service that detects process failures
- Answers queries about status of a process

Slide 11

Reliable:

- *Failed* – crashed
- *Unsuspected* – hint

Unreliable:

- *Suspected* – may still be alive
- *Unsuspected* – hint

Synchronous systems:

- Timeout
- Failure detector sends probes to detect crash failures

Asynchronous systems:

Slide 12

- ✗ Timeout gives no guarantees
 - Failure detector can track *suspected* failures
 - Combine results from multiple detectors
 - ✗ How to distinguish communication failure from process failure?
 - Ignore messages from suspected processes
 - ✓ Turn an asynchronous system into a synchronous one
-

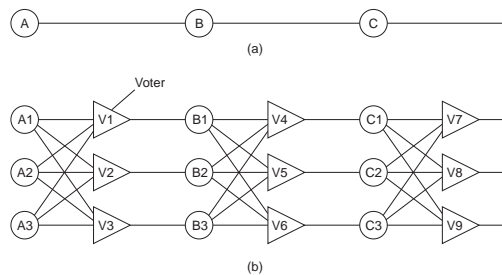
FAILURE MASKING

Try to hide occurrence of failures from other processes

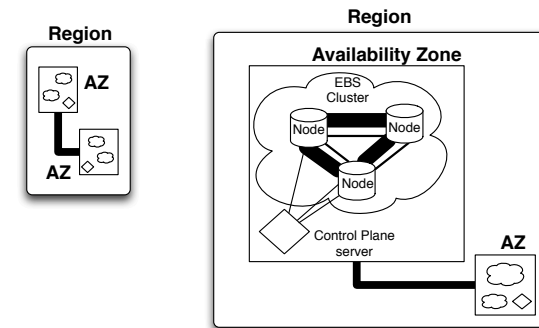
Redundancy:

- Information redundancy
- Time redundancy
- Physical redundancy

Slide 13



Slide 15



CASE STUDY: AWS FAILURE 2011

- April 21, 2011
- EBS (Elastic Block Store) in US East region unavailable for about 2 days
- 13% of volumes in one *availability zone* got stuck
- led to control API errors and outage in whole region
- led to problems with EC2 instances and RDS in most popular region
- due to reconfig error and *re-mirroring storm*.
- <http://aws.amazon.com/message/65648/>

Slide 14

RELIABLE COMMUNICATION


Slide 16

- Communication channel experiences failure
- Focus on masking crash (lost/broken connections) and omission (lost messages) failures

Slide 17

Two Army Problem:

Non-faulty processes but lossy communication.



- 1 → 2 attack!
- 2 → 1 ack
- 2: did 1 get my ack?
- 1 → 2 ack ack
- 1: did 2 get my ack ack?
- etc.

Consensus with lossy communication is impossible.
Why does TCP work?

RELIABLE POINT-TO-POINT COMMUNICATION

Slide 18

- Reliable transport protocol (e.g., TCP)
 - ✓ Masks omission failure
 - ✗ Not crash failure

Slide 19

Example: Failure and RPC:

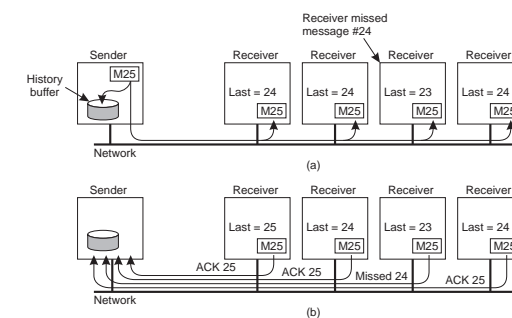
Possible failures:

- Client cannot locate server
- Request message to server is lost
- Server crashes after receiving a request
- Reply message from server is lost
- Client crashes after sending a request

How to deal with the various kinds of failure?

RELIABLE GROUP COMMUNICATION

Slide 20



SCALABILITY OF RELIABLE MULTICAST


Feedback Implosion: sender is swamped with feedback messages

Nonhierarchical Multicast:

- Use NACKs
- Feedback suppression: NACKs multicast to everyone
- Prevents other receivers from sending NACKs if they've already seen one.
- ✓ Reduces (N)ACK load on server
- ✗ Receivers have to be coordinated so they don't all multicast NACKs at same time
- ✗ Multicasting feedback also interrupts processes that successfully received message

Slide 21

Hierarchical Multicast:



Slide 22

PROCESS RESILIENCE

Protection against process failures

Groups:

- Organise identical processes into groups
- Process groups are dynamic
- Processes can be members of multiple groups
- Mechanisms for managing groups and group membership
- Deal with all processes in a group as a single abstraction

Slide 23

Flat vs Hierarchical Groups:

- Flat group: all decisions made collectively
- Hierarchical group: coordinator makes decisions

REPLICATION

Create groups using replication

Primary-Based:

- Primary-backup
- Hierarchical group
- If primary crashes others elect a new primary

Slide 24

Replicated-Write:

- Active replication or Quorum
- Flat group
- Ordering of requests (atomic multicast problem)

k Fault Tolerance:

- can survive faults in k components and still meet its specifications
- $k + 1$ replicas enough if fail-silent (or fail-stop)
- $2k + 1$ required if if byzantine

STATE MACHINE REPLICATION

Each replica executes as a state machine:

- $state + input \rightarrow output + new\ state$
- All replicas process same input in same order
- Deterministic: All **correct** replicas produce same output
- Output from incorrect replicas deviates

Input Messages:

- All replicas agree on content of input messages
- All replicas agree on order of input messages
- **Consensus** (also called **Agreement**)

Slide 25

ATOMIC MULTICAST

A message is delivered to either all processes, or none

Requires agreement about group membership


Slide 26 Process Group:

- Group view: view of the group (list of processes) sender had when message sent
- Each message uniquely associated with a group
- All processes in group have the same view

View Synchrony:

A message sent by a crashing sender is either delivered to all remaining processes (crashed after sending) or to none (crashed before sending).

Slide 27




→ view changes and messages are delivered in total order

Implementing View Synchrony:

stable message: a message that has been received by all members of the group it was sent to.

- Implemented using reliable point-to-point communication (TCP)
- Failure during multicast → only some messages delivered

Slide 28



AGREEMENT

Examples: Election, transaction commit/abort, dividing tasks among workers, mutual exclusion

Slide 29

- Previous algorithms assumed no faults
- What happens when processes can fail?
- What happens when communication can fail?
- What happens when byzantine failures are possible

We want all nonfaulty processes to reach and establish agreement (within a finite number of steps)

VARIANTS OF THE AGREEMENT PROBLEM

Consensus:

- each process proposes a value
- communicate with each other...
- all processes decide on same value
- for example, the maximum of all the proposed values

Slide 30

Interactive Consistency:

- all processes agree on a decision *vector*
- for example, the value that each of the processes proposed

Byzantine Generals:

- commander proposes a value
 - all other processes agree on the commander's value
-

Correctness of agreement:

Termination all processes eventually decide

Agreement all processes decide on the same value

Slide 31

Validity C the decided value was proposed by one of the processes

IC the decided value is a vector that reflects each of the processes proposed values

BG the decided value was proposed by the commander

CONSENSUS IN A SYNCHRONOUS SYSTEM

Slide 32

Assume:


- Execution in rounds
 - Timeout to detect lost messages
-

Byzantine Generals Problem:

Reliable communication but faulty processes.

- n generals (processes)
- m are traitors (will send incorrect and contradictory info)
- Need to know everyone else's troop strength g_i
- Each process has a vector: $\langle g_1, \dots, g_n \rangle$
- (Note: this is actually interactive consistency)

Slide 33




Byzantine agreement with Signatures:

- Digitally sign messages
- Cannot lie about what someone else said
- Avoids the impossibility result
- Can have agreement with 3 processes and 1 faulty

Slide 35

Byzantine Generals Impossibility:

Slide 34



- If m faulty processes then $2m + 1$ nonfaulty processes required for correct functioning

CONSENSUS IN AN ASYNCHRONOUS SYSTEM

Slide 36

Assume:

- Arbitrary execution time (no rounds)
- Arbitrary message delays (can't rely on timeout)

IMPOSSIBILITY OF CONSENSUS WITH ONE FAILURE

Impossible to guarantee consensus with ≥ 1 faulty process

Proof Outline:

- Slide 37**
- Fischer, Lynch, Patterson (FLP) 1985
 - the basic idea is to show circumstances under which the protocol remains forever indecisive
 - bivalent vs univalent states
 - 1. There is always a bivalent start state
 - 2. Always possible to reach a bivalent state by delaying messages
 - no termination

In practice we can get close enough

PAXOS

Goal: a collection of processes chooses a single proposed value in the presence of failure

Proposer proposes value to choose (leader)

Acceptor accept or reject proposed values

- Slide 38**
- Learner** any process interested in the result (chosen value) of the consensus

Chosen Value: value accepted by majority of acceptors

Properties:

- Only proposed values can be learned
- At most one value can be learned
- If a value has been proposed then eventually a value will be learned

PAXOS ALGORITHM: 3 PHASES

Phase 1: Propose:

- ① Propose: send a proposal $\langle \text{seq}, \text{value} \rangle$ to $\geq N/2$ acceptors
- ② Promise: acceptors reply.
 - reject if $\text{seq} < \text{seq}$ of previously accepted value
 - else accept (include last accepted value). $\text{promised} = \text{seq}$.


Phase 2: Accept:

- ① Accept: when $\geq N/2$ accept replies, proposer sends value (as received from acceptor or arbitrary):
- ② Accepted: acceptors reply.
 - reject if $\text{seq} < \text{promised}$.
 - else accepted. Remember accepted value.

Phase 3: Learn:

- ① Propagate value to Learners when $\geq N/2$ accepted replies received.

SIMPLE CASE



FAILURES

Failure Model:

channel : lose, reorder, duplicate message

process : crash (fail-stop, fail-resume)

Slide 41


Failure Cases:

- ① Acceptor fails
- ② Proposer fails
- ③ Multiple proposers

PROPOSER FAILS

- Elect a new leader
- Continue execution
- ✓ New proposer will choose any previously accepted value


Slide 43



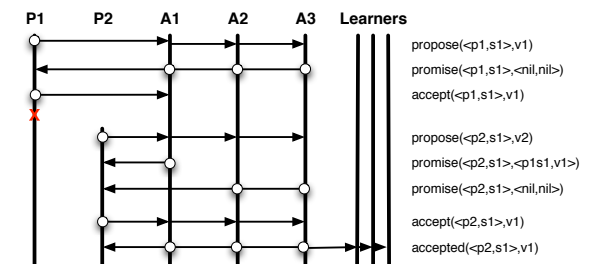
ACCEPTOR FAILS

- ✓ As long as a quorum still available
- Restart: Must remember last accepted value(s)

Slide 42



Slide 44



MULTIPLE PROPOSERS

Slide 45

- For example: crashed proposer returns and continues
- ✗ Dueling proposers
- ✗ No guaranteed termination
- ✓ Heuristics to recognise situation and back off

MULTI PAXOS

Slide 46

- Need to choose multiple values
 - agree on values
 - agree on order of values
- Run multiple *instances* of Paxos in sequence
- Each instance to choose a single value
- Add *instance id* to algorithm
- Track competed instances
- On failure, restart or join *last completed instance + 1*

USING PAXOS

Use Paxos for:

Slide 47

- Total order multicast: order messages
- State machine replication: order operations
- Leader election: choose a leader id
- Replicated storage: order writes
- View synchrony: order view changes

EXAMPLE: STATE MACHINE REPLICATION

API:

```
val run_proposer(iid, proposed_val)
run_acceptor(iid)
val learn(iid)
```

Slide 48

Client:

```
while (1){
  ...
  send(leader, nextop);
  ...
}
```

Replica: Proposer:

```
while(1) {  
  receive op  
  do {    chosen = run_proposer(i++, op);  } while (chosen != op)  
}
```

Replica: Acceptor:

Slide 49

```
while(1) {  
  run_acceptor(i++);  
}
```

Replica: Learner:

```
while(1) {  
  op = learn(i++); exec_op(op);  
}
```

OPTIMISATION AND MORE INFORMATION

Opportunities for optimisation:

- Reduce rounds
 - Phase 1: reject: return highest accepted seq
 - Phase 2: reject: return promised seq
- Reduce messages
 - Piggyback multiple requests and replies
 - Pre-propose multiple instances (assumes Proposer rarely fails)

Slide 50

More information:

Paxos Made Live - An Engineering Perspective Experiences implementing Paxos for Google's Chubby lock server. It turns out to be quite complicated.

FAILURE RECOVERY

Restoring an erroneous state to an error free state

Issues:

Slide 51

- **Reclamation of resources:**
locks, buffers held on other nodes
 - **Consistency:**
Undo partially completed operations prior to restart
 - **Efficiency:**
Avoid restarting whole system from start of computation
-

FORWARD VS. BACKWARD ERROR RECOVERY

Forward Recovery:

- Correct erroneous state without moving back to a previous state.
- Example: erasure correction - missing packet reconstructed from successfully delivered packets.
- ✗ Possible errors must be known in advance

Slide 52

Backward Recovery:

- Correct erroneous state by moving to a previously correct state
 - Example: packet retransmission when packet is lost
 - ✓ General purpose technique.
 - ✗ High overhead
 - ✗ Error can reoccur
 - ✗ Sometimes impossible to roll back (e.g. ATM has already delivered the money)
-

BACKWARD RECOVERY

General Approach:

- Restore process to *recovery point*
- Restore system by restoring all active processes

Specific Approaches:

Operation-based recovery :

- Slide 53**
- Keep *log* (or audit trail) of operations
 - Restore to recovery point by reversing changes

State-based recovery :

- Store complete state at recovery point (*checkpointing*)
- Restore process state from checkpoint (*rolling back*)

Log or checkpoint recorded on *stable* storage

Operation-Based Recovery - Logging:

Update in-place together with write-ahead logging

- Slide 54**
- Every change (update) of data is recorded in a log, which includes:
 - Data item name (for identification)
 - Old data item state (for *undo*)
 - New data item state (for *redo*)
 - Undo log is written *before* update (write-ahead log).
 - Transaction semantics
-

State-Based Recovery - Checkpointing:

Take frequent checkpoints during execution

Checkpointing:

- Slide 55**
- Pessimistic vs Optimistic
 - *Pessimistic*: assumes failure, optimised toward recovery
 - *Optimistic*: assumes infrequent failure, minimises checkpointing overhead
 - Independent vs Coordinated
 - *Coordinated*: processes synchronise to create global checkpoint
 - *Independent*: each process takes local checkpoints independently of others
 - Synchronous vs Asynchronous
 - *Synchronous*: distributed computation blocked while checkpoint taken
 - *Asynchronous*: distributed computation continues while checkpoint taken
-

Checkpointing Overhead:

- ✗ Frequent checkpointing increases overhead
- ✗ Infrequent checkpointing increases recovery cost

Decreasing Checkpointing Overhead:

Incremental checkpointing: Only write changes since last checkpoint:

- Slide 56**
- Write-protect whole address space
 - On write-fault mark page as dirty and unprotect
 - On checkpoint only write dirty pages

Asynchronous checkpointing: Use copy-on-write to checkpoint while execution continues

- Easy with UNIX `fork()`


Compress checkpoints: Reduces storage and I/O cost at the expense of CPU time

RECOVERY IN DISTRIBUTED SYSTEMS

- Failed process may have *causally affected* other processes
- Upon recovery of failed process, must undo effects on other processes
- Must roll back all affected processes
- All processes must establish recovery points
- Must roll back to a *consistent global state*

Slide 57

Domino Effect:




Slide 58

- P_1 fails → roll back: $P_1 \leadsto R_{13}$
- P_2 fails → $P_2 \leadsto R_{22}$
Orphan message m is received but not sent → $P_1 \leadsto R_{12}$
- P_3 fails → $P_3 \leadsto R_{32} \rightarrow P_2 \leadsto R_{21} \rightarrow P_1 \leadsto R_{11}, P_3 \leadsto R_{31}$

Messaging dependencies plus independent checkpointing may force system to roll back to initial state


Message Loss:



Slide 59


- Failure of $P_2 \rightarrow P_2 \leadsto R_{21}$
- Message m is now recorded as sent (by P_1) but not received (by P_2), and m will never be received after rollback
- Message m is *lost*
- Whether m is lost due to rollback or due to imperfect communication channels is indistinguishable!
- Require protocols resilient to message loss

Livelock:



$P_2 \Downarrow \rightarrow P_2 \leadsto R_{21} \rightarrow P_1 \leadsto R_{11}$. Note: n_1 in transit


Slide 60



- Pre-rollback message n_1 is received after rollback
- Forces another rollback $P_2 \leadsto R_{21}, P_1 \leadsto R_{11}$, can repeat indefinitely

CONSISTENT CHECKPOINTING


Consistent Cut:



Slide 61

Idea: collect *local checkpoints* in a coordinated way.

- Set of local checkpoints forms a *global checkpoint*.
- A global checkpoint represents a *consistent system state*.



Slide 62

- $\{R_{11}, R_{21}, R_{31}\}$ form a *strongly consistent checkpoint*:
 - No information flow during checkpoint interval
- $\{R_{12}, R_{22}, R_{32}\}$ form a *consistent checkpoint*:
 - All messages recorded as received **must be** recorded as sent

- **Strongly consistent checkpointing** requires quiescent system
 - Potentially long delays during *blocking checkpointing*
- **Consistent checkpointing** requires dealing with message loss
 - Not a bad idea anyway, as otherwise each lost message would result in a global rollback
 - Note that a consistent checkpoint may not represent an actual past system state

Slide 63

How to take a consistent checkpoint?:

- Simple solution: Each process checkpoints immediately after sending a message
- ✗ High overhead
- Reducing this to checkpointing after n messages, $n > 1$, is **not** guaranteed to produce a consistent checkpoint!
- Require some coordination during checkpointing

SYNCHRONOUS CHECKPOINTING

Processes coordinate local checkpointing so that most recent local checkpoints constitute a consistent checkpoint

Assumptions:

- Communication is via FIFO channels.
- Message loss dealt with via
 - Protocols (such as sliding window), or
 - Logging of all sent messages to stable storage
- Network will not partition

Slide 64

Local checkpoints:

permanent: part of a global checkpoint

tentative: may or may not become permanent

SYNCHRONOUS ALGORITHM

- Global checkpoint initiated by a single *coordinator*
- Based on 2PC

First Phase:

- ① Coordinator P_i takes tentative checkpoint
- ② P_i sends t message to all other processes P_j to take tentative checkpoint
- ③ P_j reply to P_i whether succeeded in taking tentative checkpoint
- ④ P_i receives *true* reply from each P_j → decides to make permanent
 P_i receives at least one *false* → decides to discard the tentative checkpoints

Slide 65

Second Phase:


- ① Coordinator P_i informs all other processes P_j of decision
- ② P_j convert or discard tentative checkpoints accordingly

Consistency ensured because no messages sent between two checkpoint messages from P_i

Slide 66

REDUNDANT CHECKPOINTS

Algorithm performs unnecessary checkpoints



Slide 67

- $\{R_{11}, R_{21}, R_{31}\}$ form a (strongly) consistent checkpoint
- Checkpoint $\{R_{12}, R_{22}, R_{32}\}$ initiated by P_1 is strongly consistent
- R_{32} is redundant, as $\{R_{12}, R_{22}, R_{31}\}$ is consistent

AVOIDING REDUNDANT CHECKPOINTS

Keep track of messages sent to avoid redundant checkpoints

- Associate each message m with *label* $m.l$, incremented at each message
- Each process maintains three vectors:
 - $last_rec_i[j] := m.l$, where m is **last** msg m received by P_i from P_j since last checkpoint ($last_rec_i[j] = 0$ if none)
 - $first_sent_i[j] := m.l$, where m is **first** msg m sent by P_i to P_j since last checkpoint ($first_sent_i[j] = 0$ if none)
 - $cohort_i := \{j | last_rec_i[j] > 0\}$, set of processes from which P_i has received a message since last checkpoint
- P_j only needs to take a checkpoint after receiving a control message t ("take tentative checkpoint") from i if $last_rec_i[j] \geq first_sent_j[i] > 0$

Slide 68

CHECKPOINTING ALGORITHM

Messages:

→ t : take tentative, p : make permanent, u : undo checkpoint

Slide 69 Initialisation: Each P sets $OK := true$, $first_sent = \{0, 0, \dots, 0\}$

Coordinator, P_i :

- ① send($t, i, last_rec_i[j]$) to all $P_j \in cohort_i$
- ② if all replies are $true$, send(p) to all $P_j \in cohort_i$
 else send(u) to all $P_j \in cohort_i$

Others, P_j : upon receiving ($t, i, last_rec_i[j]$)

- ① if OK_j **and** $last_rec_i[j] \geq first_sent_j[i] > 0$
- ② take tentative checkpoint
- ③ send($t, j, last_rec_j[k]$) to all $P_k \in cohort_j$
- ④ if all replies are $true$, $OK := true$ **else** $OK := false$
- ⑤ send(OK, j) to i

Slide 70

Others, P_j : upon receiving message $x \in \{p, u\}$ from P_i :

- ① if $x = p$ make permanent **else** discard tentative checkpoint
- ② send(x, j) to all $P_k \in cohort_j$

Note: $O(n^2)$ messages

ROLLBACK RECOVERY

First Phase:

- ① Coordinator sends “ r ” messages to all other processes to ask them to roll back
- ② Each process replies $true$, unless already in checkpoint or rollback
- ③ If all replies are $true$, coordinator decides to roll back, otherwise continue

Slide 71

Second Phase:

- ① Coordinator sends decision to other processes
- ② Processes receiving this message perform corresponding action

REDUNDANT ROLLBACKS

Processes may roll back unnecessarily

→ Can be avoided by keeping track of messages received

Avoiding Redundant Rollbacks:

- Message labelling as before
- Two additional vectors:
 - $last_sent_i[j] := m.l$, where m is **last** msg m sent by i to j since last checkpoint ($last_sent_i[j] = \infty$ if none)
 - $r_cohort_i := \{j | i \text{ communicates with } j\}$
- P_i only needs to roll back after receiving message ($r, j, last_sent_j[i]$) if $last_rec_i[j] > last_sent_j[i]$

Slide 72

ROLLBACK RECOVERY ALGORITHM

Messages:

→ r : rollback request, d : do rollback, c :continue

Initialisation: Each P sets

→ $resume := true$

→ $last_rec = \{\infty, \infty, \dots, \infty\}$

→ W according to willingness to roll back

Initiator, P_i :

- ① send($r, i, last_sent_i[j]$) to all $P_j \in r_cohort_i$
- ② if all replies are *true*, send(d) to all $P_j \in cohort_i$
 else send(c) to all $P_j \in cohort_i$

Slide 73

Others, P_j : upon receiving ($r, i, last_sent_i[j]$)

- ① if W_j **and** $last_rec_j[i] > last_sent_i[j]$ **and** $resume_j$:
- ② $resume_j = false$,
- ③ send($r, j, last_sent_j[k]$) to all $P_k \in r_cohort_j$,
- ④ if all replies are *true*, $W_j := true$ **else** $W_j := false$,
- ⑤ send(W_j, j) to i .

Slide 74

Others, P_j : upon receiving message $x \in \{c, d\}$ from P_i :

- ① if $x = d$ roll back **else** continue,
- ② send(x, j) to all $P_k \in r_cohort_j$.

ASYNCHRONOUS CHECKPOINTING

Let processes checkpoint independently (unsynchronised) and construct a consistent state during recovery.

- Source of inconsistencies are *orphan messages*.
- Consistent state can be obtained by:
 - ① Restarting failed process from latest checkpoint, and
 - ② *rolling forward* the restarted process past the point where the last message was sent prior to failure
- All send operations during roll-forward are suppressed.
- Except for timing, the result is indistinguishable from restarting from a (non-existent) checkpoint taken after the last send.
- Works as long as no message was lost.

Slide 75

MESSAGE LOGGING

- Suppressing outgoing messages during roll-forward requires knowledge of the number of messages the failed process had sent prior to failure.
 - Log the send count in *stable storage*.
- Any attempted receive of a lost message will terminate roll-forward.
 - Log all *incoming* messages in stable storage
 - During roll-forward replay incoming messages from log

Slide 76

Slide 77

Problems with Message Logging:

Roll-forward assumes *deterministic behaviour* of all processes

- Possible dependence of behaviour on uncontrollable factors (resident set).
- Possible inconsistencies between process IDs (different ID after restart!)
- Interrupt processing imposes time constraint, interrupts are *asynchronous* wrt. program messages.
 - May need to checkpoint before handling any interrupt!
- Multithreaded processes introduce a degree of non-determinism.

Require careful implementation and appropriate OS support.

OPTIMISTIC CHECKPOINTING

Asynchronous Logging:

Slide 78

- Log messages to volatile memory
- Flush to stable store asynchronously
- On failure, unflushed log is lost, resulting in inconsistent state
- Construct consistent state by rolling back orphan processes

Do synchronisation required between checkpointing and logging

Assumptions:

Communication Channels

- Reliable
- Ordered (FIFO)
- Infinite buffer size
- Finite (but arbitrary) delays

Slide 79

Event-Driven Computation

- Event is receipt of message
- Process waits for message
- Upon receipt sends 0 or more messages
- Each event logged to volatile storage as $\{e, m, msgs_sent\}$, where $msgs_sent$ is set of messages sent
- Log is flushed to stable storage asynchronously.

OPTIMISTIC CHECKPOINT RECOVERY

Slide 80

- Each process P_i keeps track of:
 - $n_rcvd_{i \leftarrow j}(E)$: # messages received from P_j (up to event E),
 - $n_sent_{i \rightarrow j}(E)$: # messages sent to P_j (up to event E).
- Upon restart, compare local message count with that of neighbours:
 - neighbour P_j is orphan if $n_rcvd_{j \leftarrow i} > n_sent_{i \rightarrow j}$,
 - must roll back until $n_rcvd_{j \leftarrow i} \leq n_sent_{i \rightarrow j}$.
- P_j 's rollback may orphan other processes (→ domino rollback).

