



# **CSE 803 Introduction to Distributed Computing**



# Naming & Cloud Computing



- ① Basic Concepts
- ② Naming Services
- ③ Attribute-based Naming (aka Directory Services)
- ④ Distributed hash tables

# WHAT IS NAMING?

Systems manage a wide collection of **entities** of different kinds. They are identified by different kinds of **names**:

→ Files (/boot/vmlinuz), Processes (1, 14293), Users (chak, ikuz, cs9243), Hosts (weill, facebook.com), . . .

**Examples of naming in distributed systems? What's the difficulty?**



# BASIC CONCEPTS

## Name:

- String of bits or characters
- Refers to an entity

## Entity:

- Resource, process, user, etc.
- **Operations** performed on entities at **access points**

## Address:

- Access point named by an **address**
- Entity address = address of entity's access point
- Multiple access points per entity
- Entity's access points may change

# BASIC CONCEPTS

➔ Name that *uniquely* identifies entity

➔ Properties:

① Refers to at most one entity

② Entity referred to by at most one identifier

③ Always refers to same entity (i.e. no reuse)

➔ Allows easy comparison of references



# SYSTEM-ORIENTED VS HUMAN-ORIENTED NAMES

## System-Oriented Names:

- Represented in machine readable form (32 or 64 bit strings)
- Structured or unstructured
- ✓ Easy to store, manipulate, compare
- ✗ Not easy to remember, hard for humans to use
- Example: inode (0x00245dad)

## Human-Oriented Names:

- Variable length character strings
- Usually structured
- Often many human-oriented names map onto a single system-oriented name
- ✓ Easy to remember and distinguish between
- ✗ Hard for machine to process
- Example: URL (<http://www.cse.unsw.edu.au/~cs9243/lectures>)

# NAME SPACES

Container for a set of related names

## Structure options:

- Flat (only leaf nodes)
- Hierarchical (Strictly hierarchical, DAG, Multiple root nodes)
- Tag-based

## Path Names (in hierarchies):

- Sequence of edge labels
- **Absolute**: if first node in path name is a root node
- **Relative**: otherwise

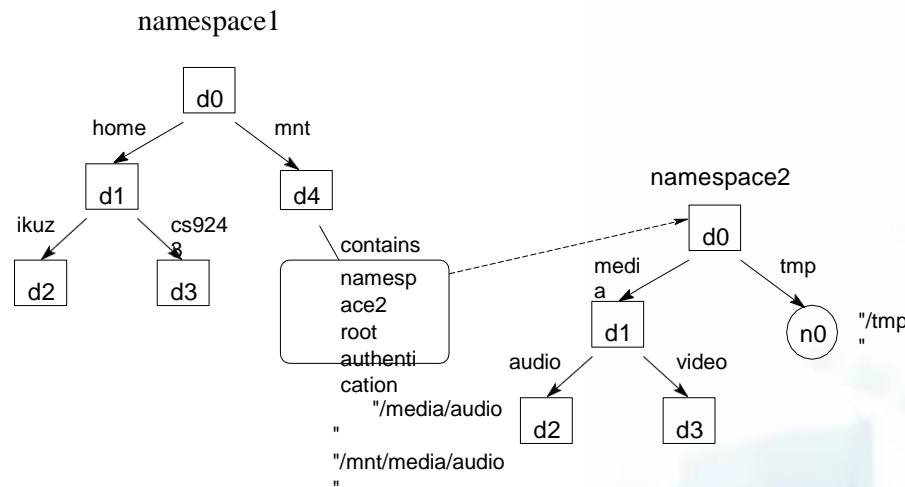
## Aliasing:

- **Alias**: another name for an entity
- **Hard link**: two or more paths to an entity in the graph
- **Soft link**: leaf node stores a (absolute) path name to another node

# NAME SPACES

## → Mounting

- Directory node stores info about a directory node in other name space
- Need: protocol, server, path name, authentication and authorisation info, keys for secure communication, etc.



## → Combining name spaces

- <http://www.cse.unsw.edu.au/~cs9243/naming-slides.pdf>
- Name Spaces: Protocol, DNS, File System



# NAMING SERVICES

A naming service provides a name space

## Name Server:

- Naming service implemented by name servers
- Implements naming service operations

## Operations:

- Lookup: resolve a path name, or element of a path name
- Add: add a directory or leaf node
- Remove: remove a subtree or leaf node
- Modify: modify the contents of a directory or leaf node

## Client:

- Invokes naming service operations

Centralised vs Distributed Naming Service

# NAME RESOLUTION

**The process of looking up a name**

## **Resolution:**

- ➔ Mapping a name onto the node referred to by the name
- ➔ Interested in the data stored by the node

## **Path Name Resolution:**

- ➔ Starts at a begin node (first element of the path name)
- Root node for absolute name
- Directory node for relative name
- ➔ Ends with data from (or a reference to) the last node (last element of path name)

## **Resolver:**

- ➔ Does name resolution on behalf of client
- ➔ In client process, in client's kernel, process on client's machine

# PARTITIONING

**Split name space over multiple servers**

## **Structured Partitioning:**

- ➔ split name space according to graph structure
- ➔ Name resolution can use zone hints to quickly find appropriate server
- ✓ Improved lookup performance due to knowledge of structure
- ✗ Rigid structure



## **Structure-free Partitioning:**

- ➔ content placed on servers independent of name space
- ✓ Flexible
- ✗ Decreased lookup performance, increased load on root



# PARTITIONING

## PARTITIONING



# DNS (DOMAIN NAME SYSTEM)

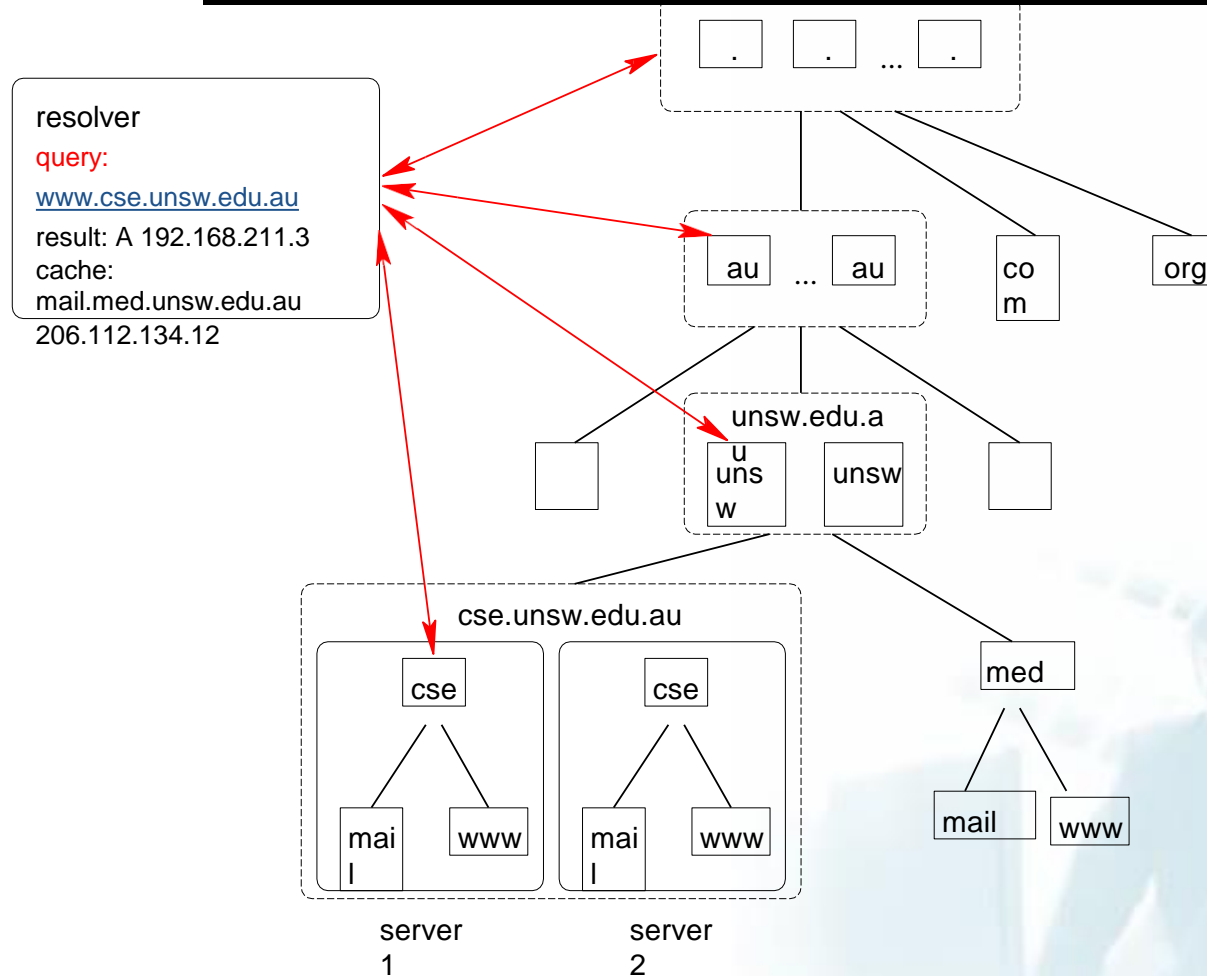
## Structure:

- Hierarchical structure (tree)
- Top-level domains (TLD) (.com, .org, .net, .au, .nl, ...)
- Zone: a (group of) directory node
- Resource records: contents of a node
- Domain: a subtree of the global tree
- Domain name: an absolute path name

Type of record	Associated entity	Description
SOA	Zone	Holds information on the represented zone
A	Host	Contains an IP address of the host this node represents
MX	Domain	Refers to a mail server to handle mail addressed to this node
SRV	Domain	Refers to a server handling a specific service
NS	Zone	Refers to a name server that implements the represented zone
CNAME	Node	Symbolic link with the primary name of the represented node
PTR	Host	Contains the canonical name of a host
HINFO	Host	Holds information on the host this node represents
TXT	Any kind	Contains any entity-specific information considered useful



# DNS (DOMAIN NAME SYSTEM)



# DNS (DOMAIN NAME SYSTEM)

- Each zone implemented by a name server

## Replication:

- Each zone replicated on at least two servers
- Updates performed on *primary*
- Contents transferred to *secondary* using *zone transfer*
- Higher levels have many more replicas (13 root servers: A-M.root-servers.net. Actually 386 replicas using anycast)

## Caching:

- Servers cache results of queries
- Original entries have time-to-live field (TTL)
- Cached data is non-authoritative, provided until TTL expires

## Name Resolution:

- Query sent to local server
- If cannot resolve locally then sent to root
- Resolved recursively or iteratively



- ① What is Cloud Computing?
- ② X as a Service
- ③ Key Challenges
- ④ Developing for the Cloud






## WHAT IS CLOUD COMPUTING?

**A style of computing in which dynamically scalable and often virtualized resources are provided as a service over the Internet. [Wikipedia]**





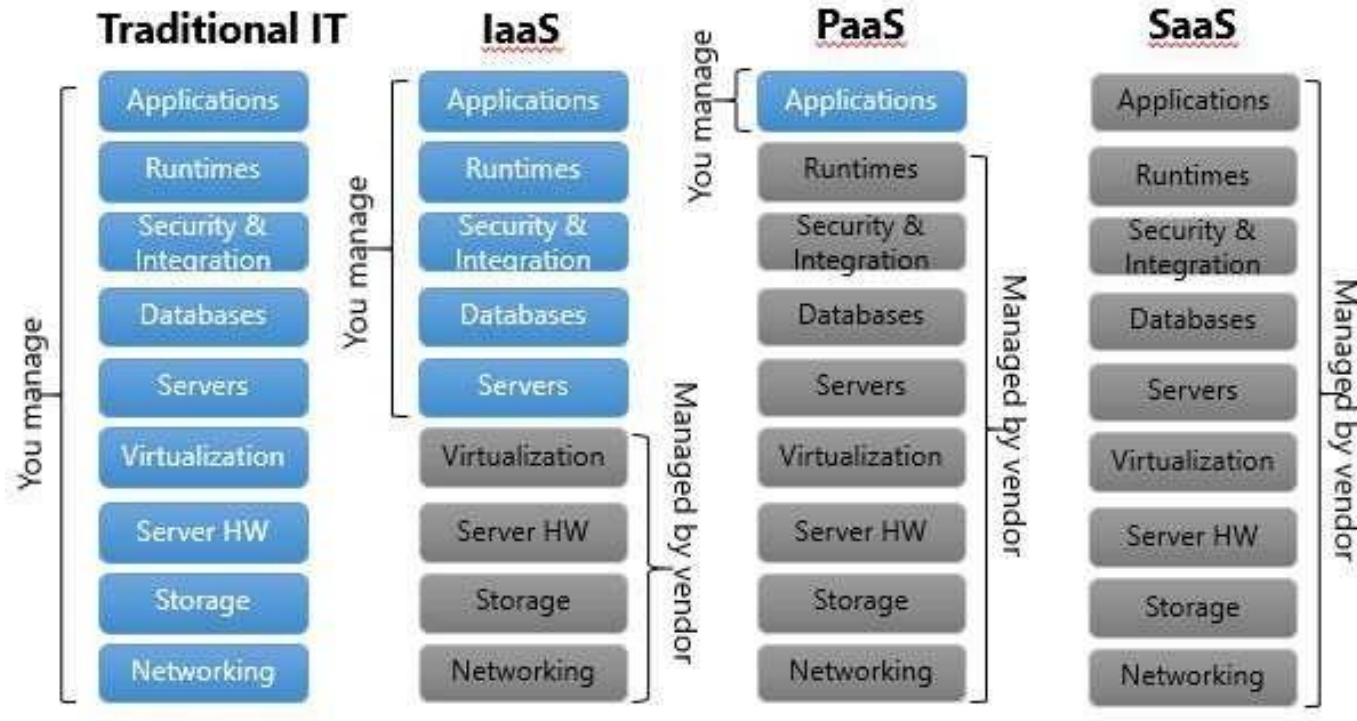
# WHAT IS CLOUD COMPUTING?



## Why is it called *Cloud*?

- services provided on virtualised resources
- virtual machines spawned on demand
- location of services no longer certain
- similar to *network cloud*

# WHAT IS CLOUD COMPUTING?



<http://www.mazikglobal.com/blog/cloud-computing-stack-saas-paas-iaas/>



# WHAT IS CLOUD COMPUTING?

Technology exposed to customers



Commercial providers



Figure from Hiroshi Wada

# KEY CHARACTERISTICS OF CLOUD COMPUTING

## KEY CHARACTERISTICS OF CLOUD COMPUTING

### SP 800-145. The NIST Definition of Cloud Computing:

#### ① **On-demand, self-service**

- get resources (CPU, storage, bandwidth etc),
- automated: as needed, right now!

#### ② **Network access**

- services accessible over the network, standard protocols

#### ③ **Pooled resources**

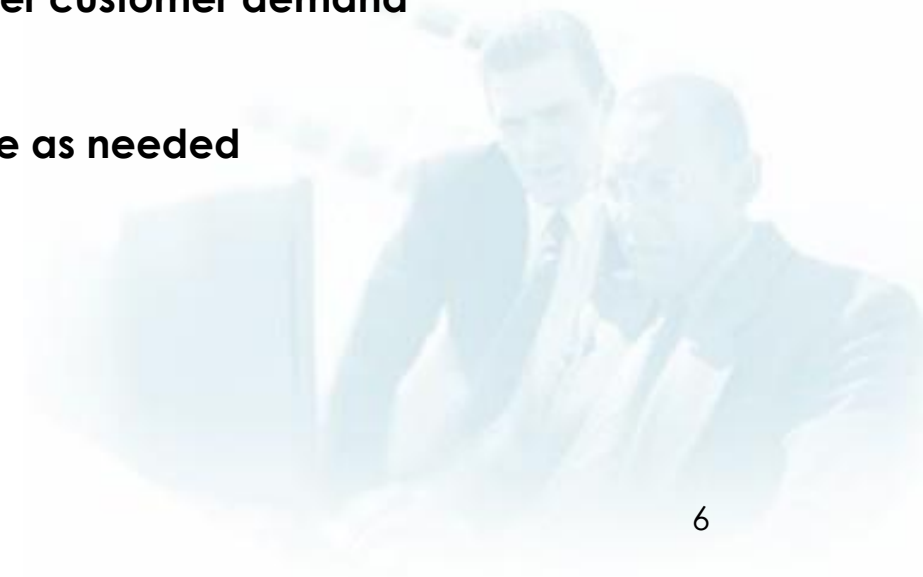
- provider: multi-tenant pool of resources
- dynamically assigned and reassigned per customer demand

#### ④ **Elasticity**

- Scalability: rapidly adjust resource usage as needed

#### ⑤ **Measured service**

- monitor resource usage
- billing for resources used



# BENEFITS

## Flexibility:

- Flexible provisioning
- Add machines on demand
- Add storage on demand

## Effort:

- Low barrier to entry
- Initial effort: no need to spec and set up physical infrastructure
- Continuing effort: no need to maintain physical infrastructure



## Speed:

- Most cloud computing services are provided self service and on demand, so even vast amounts of computing resources can be provisioned in minutes, typically with just a few mouse clicks, giving businesses a lot of flexibility and taking the pressure off capacity planning

# BENEFITS

## Cost:

- Cloud computing eliminates the capital expense of buying hardware and software
- setting up and running on-site datacenters—the racks of servers, the round-the-clock electricity for power and cooling, and the IT experts for managing the infrastructure. It adds up fast.
- Low initial capital expenditure
- Avoid costs of over-provisioning for scalability
- Pay for what you use



in "Developing and Extending Applications for Windows Azure with Visual Studio"

# BENEFITS

## Global scale

→ The benefits of cloud computing services include the ability to scale elastically. In cloud speak, that means delivering the right amount of IT resources—for example, more or less computing power, storage, bandwidth—right when they're needed, and from the right geographic location.

## Security

→ Many cloud providers offer a broad set of policies, technologies, and controls that strengthen your security posture overall, helping protect your data, apps, and infrastructure from potential threats.

- Redundancy
- Trust reliability of provider
- Data backups
- *What happens when provider goes down?*
- *What about Security? Privacy?*





# BENEFITS

## **Performance:**

→ The biggest cloud computing services run on a worldwide network of secure datacenters, which are regularly upgraded to the latest generation of fast and efficient computing hardware. This offers several benefits over a single corporate datacenter, including reduced network latency for applications and greater economies of scale.

## **Reliability**

→ Cloud computing makes data backup, disaster recovery, and business continuity easier and less expensive because data can be mirrored at multiple redundant sites on the cloud provider's network.


## **Productivity**

→ On-site datacenters typically require a lot of “racking and stacking”—hardware setup, software patching, and other time-consuming IT management chores. Cloud computing removes the need for many of these tasks, so IT teams can spend time on achieving more important business goals.

## **Speed**

→ Most cloud computing services are provided self service and on demand, so even vast amounts of computing resources can be provisioned in minutes, typically with just a few mouse clicks, giving businesses a lot of flexibility and taking the pressure off capacity planning.

# Types of cloud computing



**Public:** open services available to everyone

**Private:** owned, operated, and available to specific organisation **Is this still cloud computing?**

**Hybrid:** system uses some private cloud services and some public cloud services.

<http://blog.nskinc.com/IT-Services-Boston/bid/32590/Private-Cloud-or-Public-Cloud>

# Types of cloud services

Most cloud computing services fall into four broad categories: infrastructure as a service (IaaS), platform as a service (PaaS), serverless, and software as a service (SaaS). These are sometimes called the cloud computing "stack" because they build on top of one another. Knowing what they are and how they're different makes it easier to accomplish your business goals.

## **Infrastructure as a service (IaaS)**

The most basic category of cloud computing services. With IaaS, you rent IT infrastructure—servers and virtual machines (VMs), storage, networks, operating systems—from a cloud provider on a pay-as-you-go basis.

## **Platform as a service (PaaS)**

Platform as a service refers to cloud computing services that supply an on-demand environment for developing, testing, delivering, and managing software applications. PaaS is designed to make it easier for developers to quickly create web or mobile apps, without worrying about setting up or managing the underlying infrastructure of servers, storage, network, and databases needed for development.

## **Serverless computing**

Overlapping with PaaS, serverless computing focuses on building app functionality without spending time continually managing the servers and infrastructure required to do so. The cloud provider handles the setup, capacity planning, and server management for you. Serverless architectures are highly scalable and event-driven, only using resources when a specific function or trigger occurs.

# Types of cloud services

## **Software as a service (SaaS)**

Software as a service is a method for delivering software applications over the Internet, on demand and typically on a subscription basis. With SaaS, cloud providers host and manage the software application and underlying infrastructure, and handle any maintenance, like software upgrades and security patching. Users connect to the application over the Internet, usually with a web browser on their phone, tablet, or PC.



# INFRASTRUCTURE AS A SERVICE: IAAS

## **Service provider provides:**

- ➔ **Server and network hardware**
- ➔ **Virtual machines**
- ➔ **IP addresses**
- ➔ **Services to manage VMs (create, start, stop, migrate)**
- ➔ **Optional: storage, database, synchronization, communication**

## **Client provides:**

- ➔ **OS and OS environment**
- ➔ **Web server, DBMS, etc.**
- ➔ **Middleware**
- ➔ **Application software**



# INFRASTRUCTURE AS A SERVICE: IAAS

## Challenges

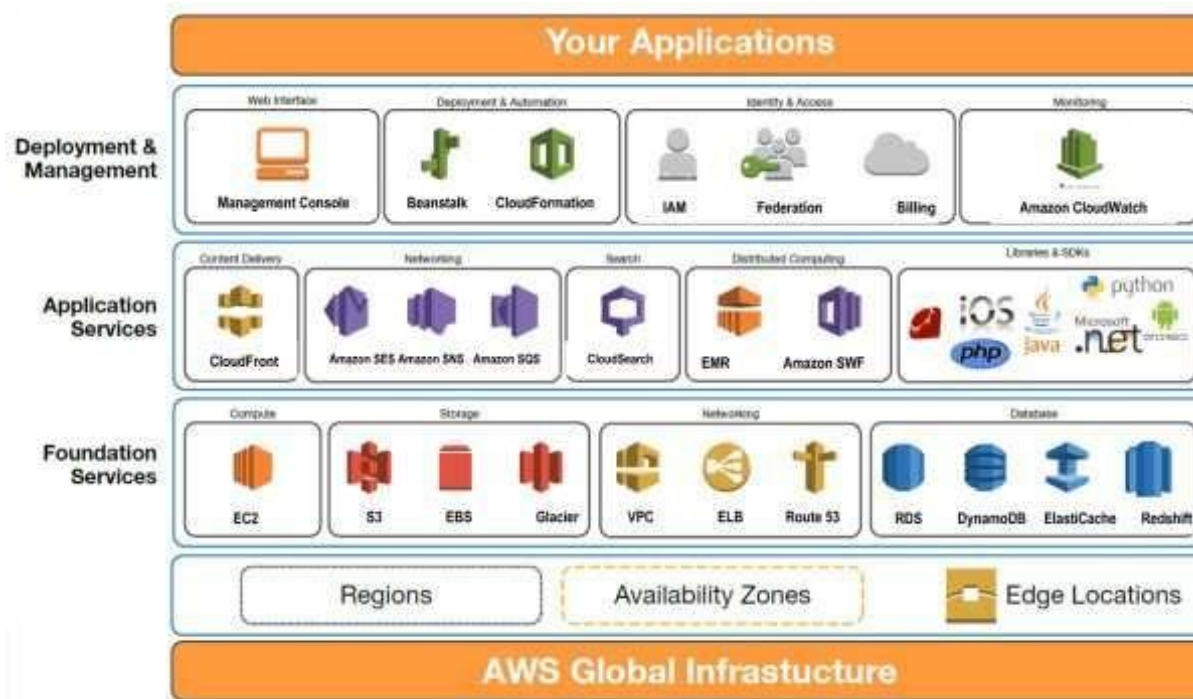
- **Transparency (naming, redirection)**
- **Scalability: replication and load balancing decisions**
- **Synchronization and coordination**
- **Security**
- **Fault tolerance**
- **Software maintenance and sys admin**

### Challenges – Provider:

- **Hardware provisioning and maintenance**
- **Load management**
- **IP address management, DNS management**
- **Infrastructure fault tolerance**
- **Monitoring, logging, billing**
- **Storage**



# EXAMPLE 1: AMAZON WEB SERVICES (AWS)



- Elastic Compute Cloud (EC2)
- Simple Storage Solution (S3)
- Simple DB
- Simple Queue Service

## EXAMPLE 1: AMAZON WEB SERVICES (AWS)

➔ **Instances: virtual cores, memory, storage**

• **instance types (cpu,memory,net, storage options):**

• **t, m, c, p, g, x, r, i, d**

• **micro, small, medium, large, xlarge, ...**

➔ **Cost:**

• **free tier: limited instances, free CPU hours**

• **on-demand: \$0.007 - \$39 per hour**

• **reserved: 1-3 years, discounted, fixed cost**

➔ **Launch Amazon Machine Image (AMI) on instances**

➔ **Preconfigured or custom images**



## USING EC2



# USING EC2

1. Grab your credit card and create an account (10 min).  
Open the EC2 Dashboard.

The screenshot shows the Amazon EC2 Dashboard for the Asia Pacific (Sydney) region. The left sidebar contains navigation links for EC2 Dashboard, Events, Tags, Reports, Limits, INSTANCES, IMAGES, ELASTIC BLOCK STORE, and NETWORK & SECURITY. The main content area displays 'Resources' (0 Running Instances, 0 Elastic IPs, 0 Dedicated Hosts, 0 Snapshots, 0 Volumes, 0 Load Balancers, 4 Key Pairs, 1 Security Groups, 0 Placement Groups) and a 'Create Instance' section with a 'Launch Instance' button. A red arrow points to the 'Launch Instance' button, and another red arrow points to the 'Asia Pacific (Sydney)' region selection in the right sidebar. A third red arrow points to the 'Launch Instance' button. A yellow callout box with the text '3. Hit this button!' is positioned over the 'Launch Instance' button.

Resources

You are using the following Amazon EC2 resources in the Asia Pacific (Sydney) region:

- 0 Running Instances
- 0 Elastic IPs
- 0 Dedicated Hosts
- 0 Snapshots
- 0 Volumes
- 0 Load Balancers
- 4 Key Pairs
- 1 Security Groups
- 0 Placement Groups

Learn more about the latest in AWS Compute from AWS re:Invent 2017 by viewing the EC2 Videos.

Create Instance

To start using Amazon EC2 you will want to launch a virtual server, known as an Amazon EC2 Instance.

**Launch Instance**

Note: Your instances will launch in the Asia Pacific (Sydney) region.

Service Health

Service Status:

- Asia Pacific (Sydney): This service is operating normally.

Availability Zone Status:

US East (N. Virginia)

US East (Ohio)

US West (N. California)

US West (Oregon)

Asia Pacific (Mumbai)

Asia Pacific (Seoul)

Asia Pacific (Singapore)

**Asia Pacific (Sydney)**

Asia Pacific (Tokyo)

Canada (Central)

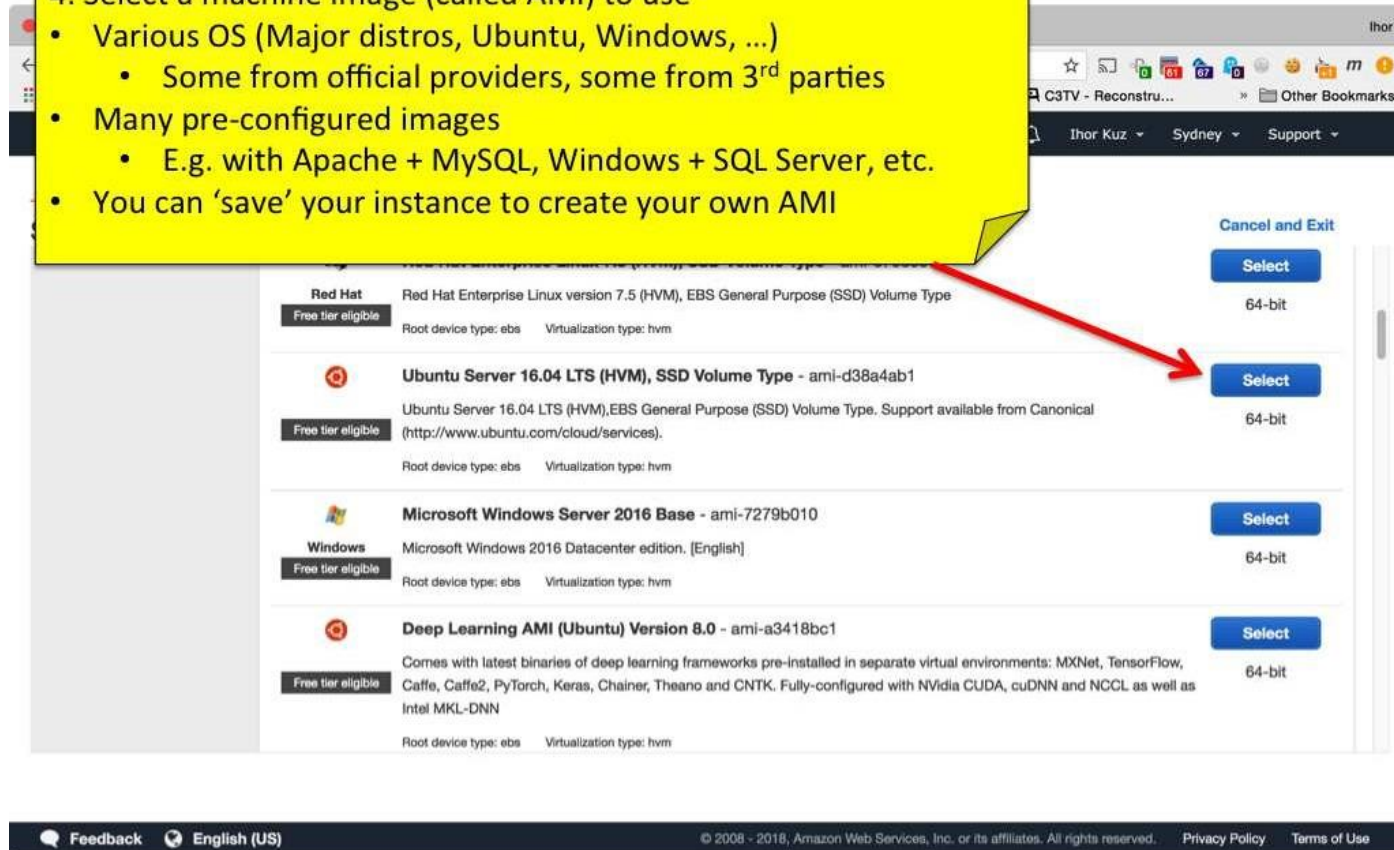
EU (Frankfurt)

3. Hit this button!

# USING EC2

## 4. Select a machine image (called AMI) to use

- Various OS (Major distros, Ubuntu, Windows, ...)
  - Some from official providers, some from 3<sup>rd</sup> parties
- Many pre-configured images
  - E.g. with Apache + MySQL, Windows + SQL Server, etc.
- You can 'save' your instance to create your own AMI



The screenshot shows the AWS Management Console interface for selecting an Amazon Machine Image (AMI). The page lists several AMIs, each with a logo, name, description, and a 'Select' button. A red arrow points from the yellow callout box to the 'Select' button for the Ubuntu Server 16.04 LTS AMI.

Logo	AMI Name	Description	Root device type	Virtualization type	Architecture
Red Hat	Red Hat Enterprise Linux version 7.5 (HVM), SSD General Purpose (SSD) Volume Type	Red Hat Enterprise Linux version 7.5 (HVM), EBS General Purpose (SSD) Volume Type	ebs	hvm	64-bit
Ubuntu	Ubuntu Server 16.04 LTS (HVM), SSD Volume Type - ami-d38a4ab1	Ubuntu Server 16.04 LTS (HVM), EBS General Purpose (SSD) Volume Type. Support available from Canonical ( <a href="http://www.ubuntu.com/cloud/services">http://www.ubuntu.com/cloud/services</a> ).	ebs	hvm	64-bit
Windows	Microsoft Windows Server 2016 Base - ami-7279b010	Microsoft Windows 2016 Datacenter edition. [English]	ebs	hvm	64-bit
Deep Learning AMI (Ubuntu)	Deep Learning AMI (Ubuntu) Version 8.0 - ami-a3418bc1	Comes with latest binaries of deep learning frameworks pre-installed in separate virtual environments: MXNet, TensorFlow, Caffe, Caffe2, PyTorch, Keras, Chainer, Theano and CNTK. Fully-configured with Nvidia CUDA, cuDNN and NCCL as well as Intel MKL-DNN	ebs	hvm	64-bit

At the bottom of the console, there is a footer with a 'Feedback' link, the language 'English (US)', and copyright information: '© 2008 - 2018, Amazon Web Services, Inc. or its affiliates. All rights reserved.' followed by links to 'Privacy Policy' and 'Terms of Use'.

# USING EC2

5. Determine the amount of resources to allocate. Price varies, e.g:

- t2.micro: USD 0.0146/hour (Linux) USD 0.0192/hour (Win)
- t2.medium: USD 0.0584/hour (Linux) USD 0.0764/hour (Win)
- m5.large: USD 0.12/hour (Linux) USD 0.212/hour (Win)

Additional costs for other software (e.g. SQL Server)

	Family	Type	vCPUs	Memory (GiB)	Instance Storage (GB)	EBS-Optimized Available	Network Performance	IPv6 Support
	General purpose	t2.nano	1	0.5	EBS only	-	Low to Moderate	Yes
	General purpose	t2.micro Free tier eligible	1	1	EBS only	-	Low to Moderate	Yes
	General purpose	t2.small	1	2	EBS only	-	Low to Moderate	Yes
<input checked="" type="checkbox"/>	General purpose	t2.medium	2	4	EBS only	-	Low to Moderate	Yes
	General purpose	t2.large	2	8	EBS only	-	Low to Moderate	Yes
	General purpose	t2.xlarge	4	16	EBS only	-	Moderate	Yes
	General purpose	t2.2xlarge	8	32	EBS only	-	Moderate	Yes
	General purpose	m5.large	2	8	EBS only	Yes	Up to 10 Gigabit	Yes
	General purpose	m5.xlarge	4	16	EBS only	Yes	Up to 10 Gigabit	Yes

Cancel

Previous

Review and Launch

Next: Configure Instance Details

Feedback

English (US)

© 2008 - 2018, Amazon Web Services, Inc. or its affiliates. All rights reserved.

Privacy Policy

Terms of Use





# USING EC2

6. Done! (< 5 minutes in total)

- Set SSH key, configure firewall, etc.
- Each machine has a randomly assigned public IP address and DNS name, e.g.:

`ec2-54-252-240-203.ap-southeast-2.compute.amazonaws.com`

EC2 Management Console

Launch Instance Connect Actions

Instance ID: i-0720818ae0d9b98ec

Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status
	i-0720818ae0d9b98ec	t2.medium	ap-southeast-2b	running	Initializing	None

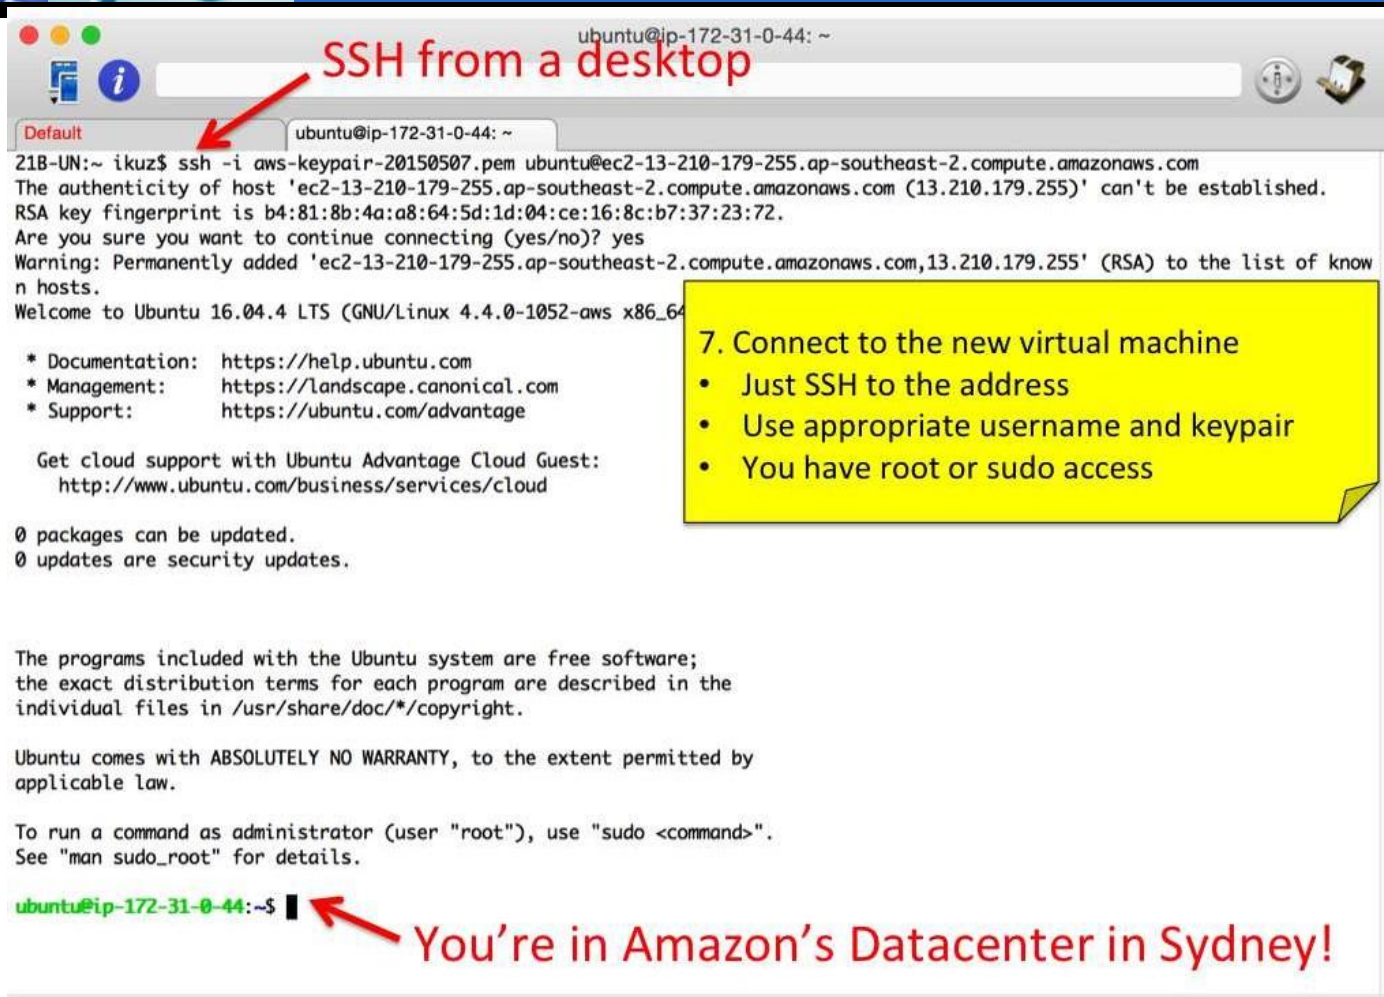
Instance: i-0720818ae0d9b98ec Public DNS: ec2-13-210-179-255.ap-southeast-2.compute.amazonaws.com

**I got my instance!**

Description	Status Checks	Monitoring	Tags
Instance ID	i-0720818ae0d9b98ec	Public DNS (IPv4)	ec2-13-210-179-255.ap-southeast-2.compute.amazonaws.com
Instance state	running	IPv4 Public IP	13.210.179.255
Instance type	t2.medium	IPv6 IPs	-
Elastic IPs		Private DNS	ip-172-31-0-44.ap-southeast-2.compute.internal
Availability zone	ap-southeast-2b	Private IPs	172.31.0.44
Security groups	launch-wizard-1 . view inbound rules	Secondary private IPs	
Scheduled events	No scheduled events	VPC ID	vpc-7a0dd91f
AMI ID	ubuntu/images/hvm-ssd/ubuntu-xenial-16.04-amd64-server-20180306 (ami-d38a4ab1)	Subnet ID	subnet-b8e14edd

Feedback English (US) © 2008 - 2018, Amazon Web Services, Inc. or its affiliates. All rights reserved. Privacy Policy Terms of Use

# USING EC2



SSH from a desktop

```
ubuntu@ip-172-31-0-44: ~  
21B-UN:~ ikuz$ ssh -i aws-keypair-20150507.pem ubuntu@ec2-13-210-179-255.ap-southeast-2.compute.amazonaws.com  
The authenticity of host 'ec2-13-210-179-255.ap-southeast-2.compute.amazonaws.com (13.210.179.255)' can't be established.  
RSA key fingerprint is b4:81:8b:4a:a8:64:5d:1d:04:ce:16:8c:b7:37:23:72.  
Are you sure you want to continue connecting (yes/no)? yes  
Warning: Permanently added 'ec2-13-210-179-255.ap-southeast-2.compute.amazonaws.com,13.210.179.255' (RSA) to the list of known hosts.  
Welcome to Ubuntu 16.04.4 LTS (GNU/Linux 4.4.0-1052-aws x86_64)  
  
* Documentation:  https://help.ubuntu.com  
* Management:    https://landscape.canonical.com  
* Support:       https://ubuntu.com/advantage  
  
Get cloud support with Ubuntu Advantage Cloud Guest:  
http://www.ubuntu.com/business/services/cloud  
  
0 packages can be updated.  
0 updates are security updates.  
  
The programs included with the Ubuntu system are free software;  
the exact distribution terms for each program are described in the  
individual files in /usr/share/doc/*/copyright.  
  
Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by  
applicable law.  
  
To run a command as administrator (user "root"), use "sudo <command>".  
See "man sudo_root" for details.  
  
ubuntu@ip-172-31-0-44:~$
```

7. Connect to the new virtual machine

- Just SSH to the address
- Use appropriate username and keypair
- You have root or sudo access

You're in Amazon's Datacenter in Sydney!

# USING EC2

If you need Windows, launch a Windows instance and remote-desktop to it.

compute.amazonaws.com

Hostname: EC2AMAZ-2U62EF7  
Instance ID: i-0c31f5425c6bb1612  
Public IP Address: 54.252.240.203  
Private IP Address: 172.31.5.245  
Instance Size: t2.micro  
Availability Zone: ap-southeast-2b  
Architecture: ARM64  
Total Memory: 1 GB  
Network Performance: Low to Moderate

File Home View  
Clipboard Select Image Tools Brushes Shapes Size Color 1 Color 2 Colors Edit colors  
Untitled - Paint  
Hello!  
768 x 480px 100%

Windows taskbar: 1:40 AM 5/15/2018

**You're in Amazon's Datacenter in Sydney!**





# USING EC2

The screenshot shows the AWS EC2 Management Console. On the left is a navigation menu with categories like INSTANCES, IMAGES, ELASTIC BLOCK STORE, and NETWORK & SECURITY. The main area displays a table of instances. Two instances are listed, both with a state of 'terminated'. Below the table, the details for instance i-0720818ae0d9b98ec are shown, including its type (t2.medium), availability zone (ap-southeast-2b), and various IP addresses.

8. Terminate (decommission) or stop (shutdown/hibernate) instances when they are not in use

- Instances cost you by time – not by actual resource usage
- Consider using a script to stop instances at a convenient time (say midnight)
- Restart instances manually when you next need them.

Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status
	i-0720818a...	t2.medium	ap-southeast-2b	terminated		None
	i-0c31f5425...	t2.micro	ap-southeast-2b	terminated		None

Instance: i-0720818ae0d9b98ec Public DNS: -

Description	Status Checks	Monitoring	Tags
Instance ID	i-0720818ae0d9b98ec		
Instance state	terminated		
Instance type	t2.medium		
Elastic IPs			
Availability zone	ap-southeast-2b		
Security groups	-		
Scheduled events	-		
Public DNS (IPv4)	-		
IPv4 Public IP	-		
IPv6 IPs	-		
Private DNS	-		
Private IPs	-		
Secondary private IPs	-		
VPC ID	-		






# USING EC2

<b>▼ Asia Pacific (Sydney)</b>			<b>\$2.33</b>
Amazon Elastic Compute Cloud running Linux/UNIX			\$0.04
\$0.0146 per On Demand Linux t2.micro Instance Hour	1.366 Hrs		\$0.02
\$0.0292 per On Demand Linux t2.small Instance Hour	0.536 Hrs		\$0.02
EBS			\$2.29
\$0.055 per GB-Month of snapshot data stored - Asia Pacific (Sydney)	5.449 GB-Mo		\$0.30
\$0.12 per GB-month of General Purpose SSD (gp2) provisioned storage - Asia Pacific (Sydney)	16.577 GB-Mo		\$1.99
<b>▼ US East (N. Virginia)</b>			<b>\$0.01</b>
EBS			\$0.01
\$0.05 per GB-Month of snapshot data stored - US East (Northern Virginia)	0.231 GB-Mo		\$0.01
<b>▼ Simple Queue Service</b>			<b>\$0.00</b>
<b>▼ Asia Pacific (Sydney)</b>			<b>\$0.00</b>
Amazon Simple Queue Service APS2-Requests-Tie			\$0.00
First 1,000,000 Amazon SQS Requests per month are free			\$0.00
<b>▼ Simple Storage Service</b>			<b>\$0.00</b>
<b>▼ No Region</b>			<b>-\$0.04</b>

9. Check the cost in near real-time

- Hours to run virtual machines
- Network in/out
- VPN
- Disk access
- # of request made to services
- ...

# RELIABILITY



<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-regions-availability-zones.html>

## Regions and Availability Zones:

- 99.95% availability per service region
- Regions: geographically dispersed, independent
- Availability zones: contained in Regions
- Availability zones: isolated from failures in other zones, but connected

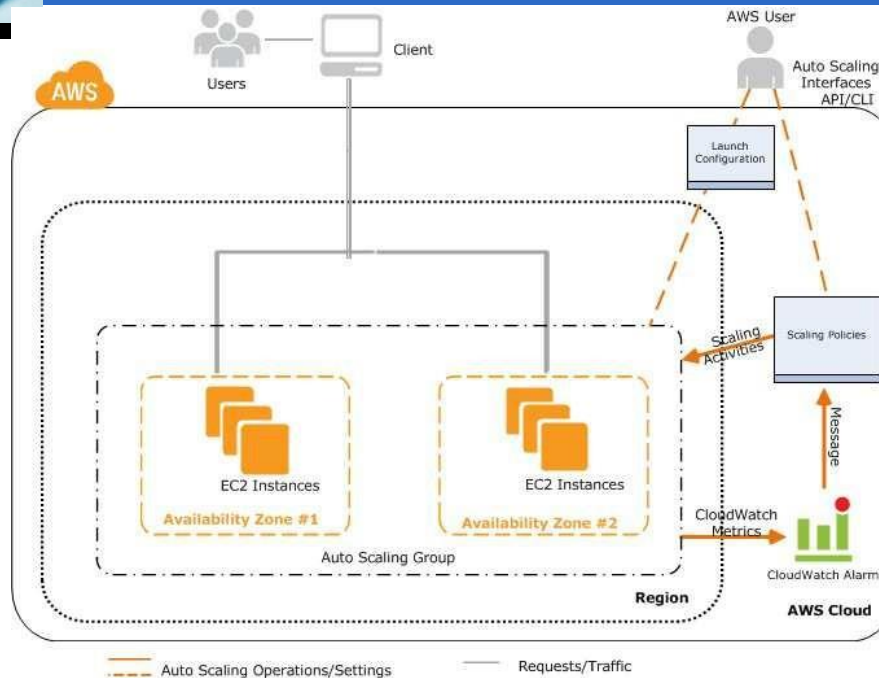
# RELIABILITY

- IP address associated with account
- Dynamic remapping to specific instances
- instance has *private IP address* and *public IP address*
- *Elastic IP* can be mapped (and re-mapped) to private IP

## Elastic Load Balancing:

- Distributes traffic across instances
- Monitors 'health' of instances: customisable
- Routes to healthy instances

# RELIABILITY



## Auto Scaling:

- Automatically start or stop new instances
- User-defined conditions
- manual (minimum group size), schedule
- instance health, CloudWatch input

<https://docs.aws.amazon.com/autoscaling/ec2/userguide/what-is-amazon-ec2-auto-scaling.html>

# RELIABILITY

## ➔ Infrastructure Security

- Data centre physical security
- Software and hardware maintenance
- Monitoring and Testing (automatic and manual)

## ➔ Application Security

- API access control (access keys)
- Firewall settings for instances (security groups)
- Virtual Private Cloud (VPC): private or public subnetworks
- Encrypted storage support
- Logging



## STORAGE

### Elastic Block Store:

- Network Attached Storage (NAS) (servers with disks)
- Block level storage volumes
- Mounted as block device (e.g. disk) on an instance
- Physical Servers and Disks shared by customers (no caching, competing for disk and net IO)
- Replicated in Availability zone
- Cost: per GB/per month

# STORAGE

- ➔ Buckets: store objects
- Can be placed in specific regions
- ➔ Objects: data and metadata
- metadata: key-value pairs describing the object
- identified by key (unique within a bucket)
- versioned
- ➔ Consistency:
  - highly replicated
  - eventual consistency, no locking
  - atomic object update
- ➔ Access control



# STORAGE

- Point in time copy of EBS volume
- Stored in S3
- Differential
- Can be used to bootstrap image

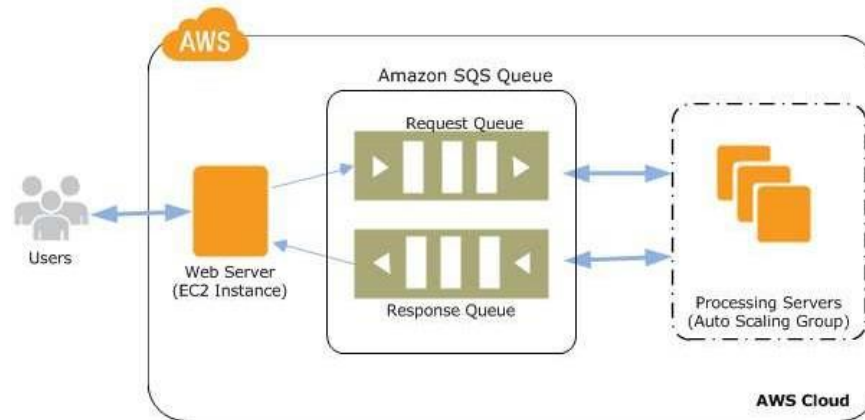
## Simple Database Service (SimpleDB):

- Non-relational database: key-value
- Partitioned into *domains*
- Consistency
  - highly replicated
  - eventual consistency
- Typical uses: logging, indexing S3 data
- Erlang!
- Replaced by DynamoDB





# COMMUNICATION



## Simple Queue Service (SQS):

- Message-queue oriented communication service
- Persistent, asynchronous messaging
- At-least once delivery guarantee
- No ordering guarantee
- Access control

<https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/>

# PLATFORM AS A SERVICE

## **Service provider provides:**

- ➔ **Hardware infrastructure**
- ➔ **OS and platform software (middleware)**
- ➔ **Distributed storage management**
- ➔ **Load balancing, replication, migration**
- ➔ **Management and Monitoring services**

## **Client provides:**

- ➔ **Application**



# PLATFORM AS A SERVICE

## Challenges

- Learn new API and environment
- Follow API
- Optimise to limits of API and platform
- Security for own app

### Challenges – Provider:

- Transparency (naming, redirection)
- Scalability: replication and load balancing decisions
- Synchronisation and coordination
- Security
- Fault tolerance
- Monitoring
- Software maintenance and sys admin

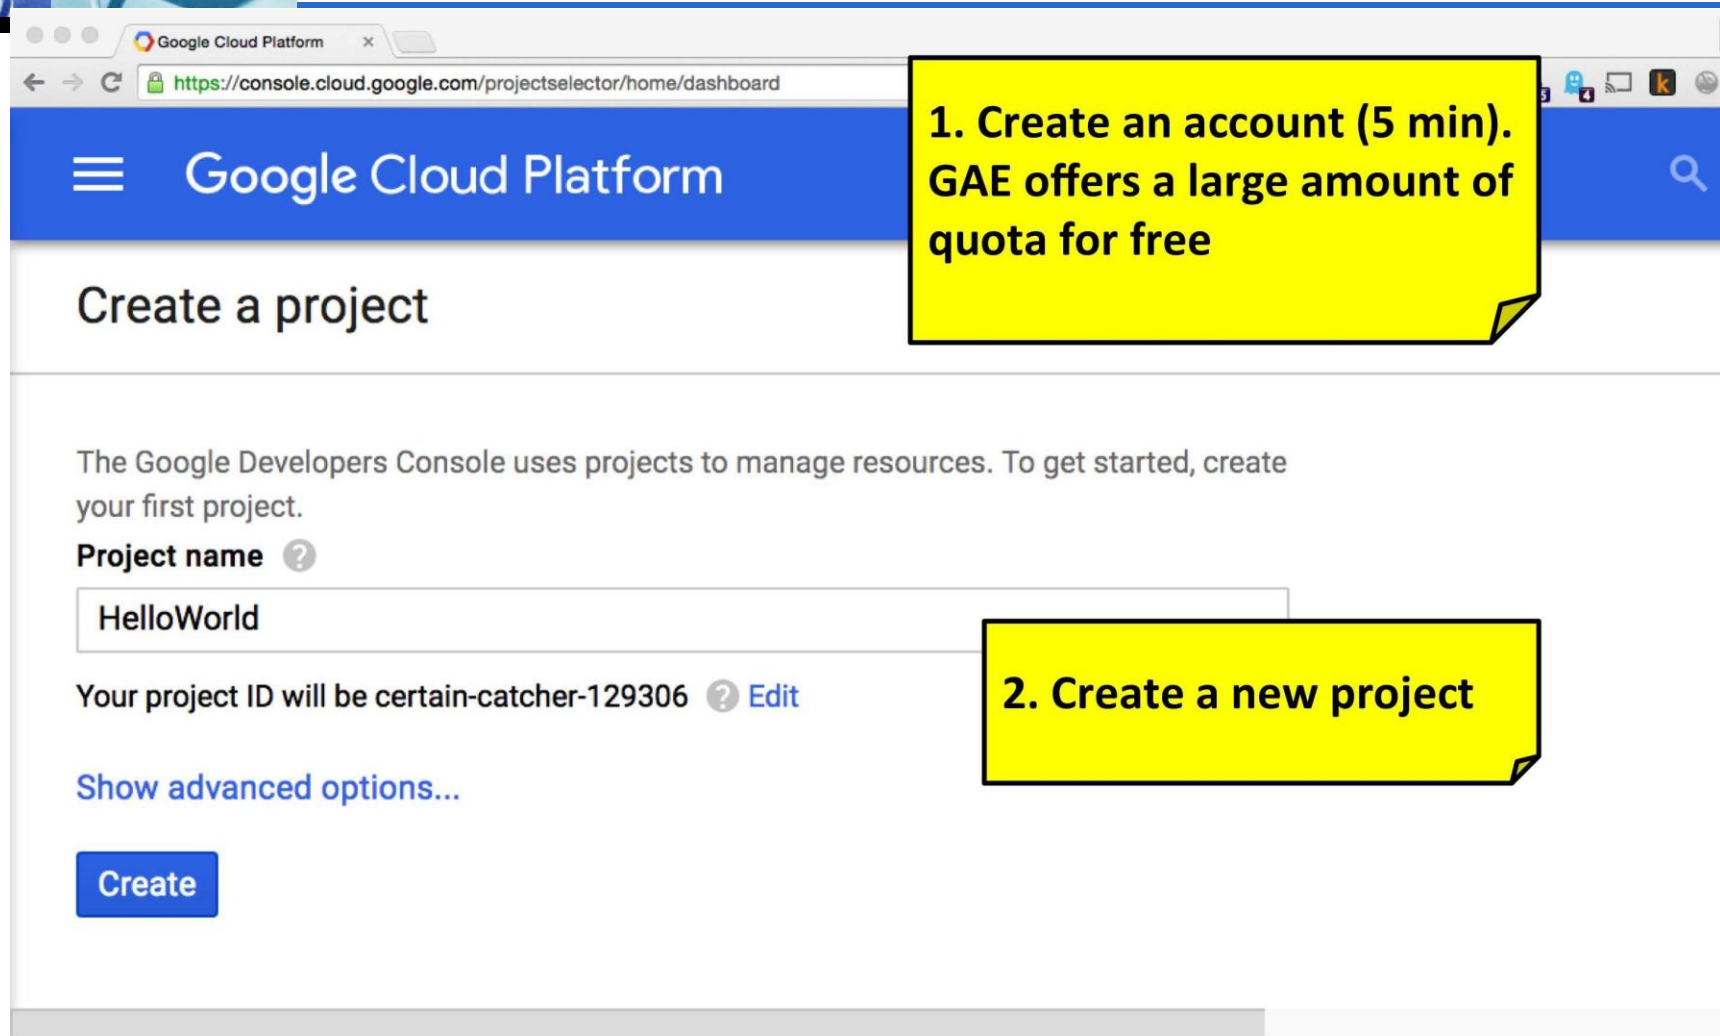


## EXAMPLE 2: APP ENGINE



- Various development languages (Python, Java, PHP, Go)
- ... and runtime environments
- Storage based on Big Table
- Optimisation via Memcache
- Lots of APIs
- Per use billing
- Transparent scaling

# EXAMPLE 2: APP ENGINE



Google Cloud Platform

<https://console.cloud.google.com/projectselector/home/dashboard>

## Create a project

The Google Developers Console uses projects to manage resources. To get started, create your first project.

**Project name** ?

HelloWorld

Your project ID will be certain-catcher-129306 ? [Edit](#)

[Show advanced options...](#)

[Create](#)

**1. Create an account (5 min).  
GAE offers a large amount of  
quota for free**

**2. Create a new project**

# EXAMPLE 2: APP ENGINE

## 3. Write an application using GAE's framework

```
le Edit Options Buffers Tools Python Help
port webapp2

class MainPage(webapp2.RequestHandler):
    def get(self):
        self.response.headers['Content-Type'] = 'text/plain'
        self.response.write('sup, World!')

p = webapp2.WSGIApplication([
    ('/', MainPage),
    debug=True])


--:---F1 helloworld.py All L1 (Python)-----
ading python...done
```

```
2-1B:helloworld ikuz$ appcfg.py -A certain-catcher-129306 -V v1 up
ate `pwd`
05:13 PM Application: certain-catcher-129306 (was: None); version: v1
(was: None)
05:13 PM Host: appengine.google.com
05:13 PM Starting update of app: certain-catcher-129306, version: v1
05:13 PM Getting current resource limits.
05:13 PM Scanning files on local disk.
05:13 PM Cloning 2 application files.
05:13 PM Compilation starting.
05:13 PM Compilation completed.
05:13 PM Starting deployment.
05:14 PM Checking if deployment succeeded.
05:14 PM Deployment successful.
05:14 PM Checking if updated app version is serving.
05:14 PM Completed update of app: certain-catcher-129306, version: v1
2-1B:helloworld ikuz$
```

## 4. Deploy your application on GAE!



# EXAMPLE 2: APP ENGINE



sup, World!

## 5. Running application.

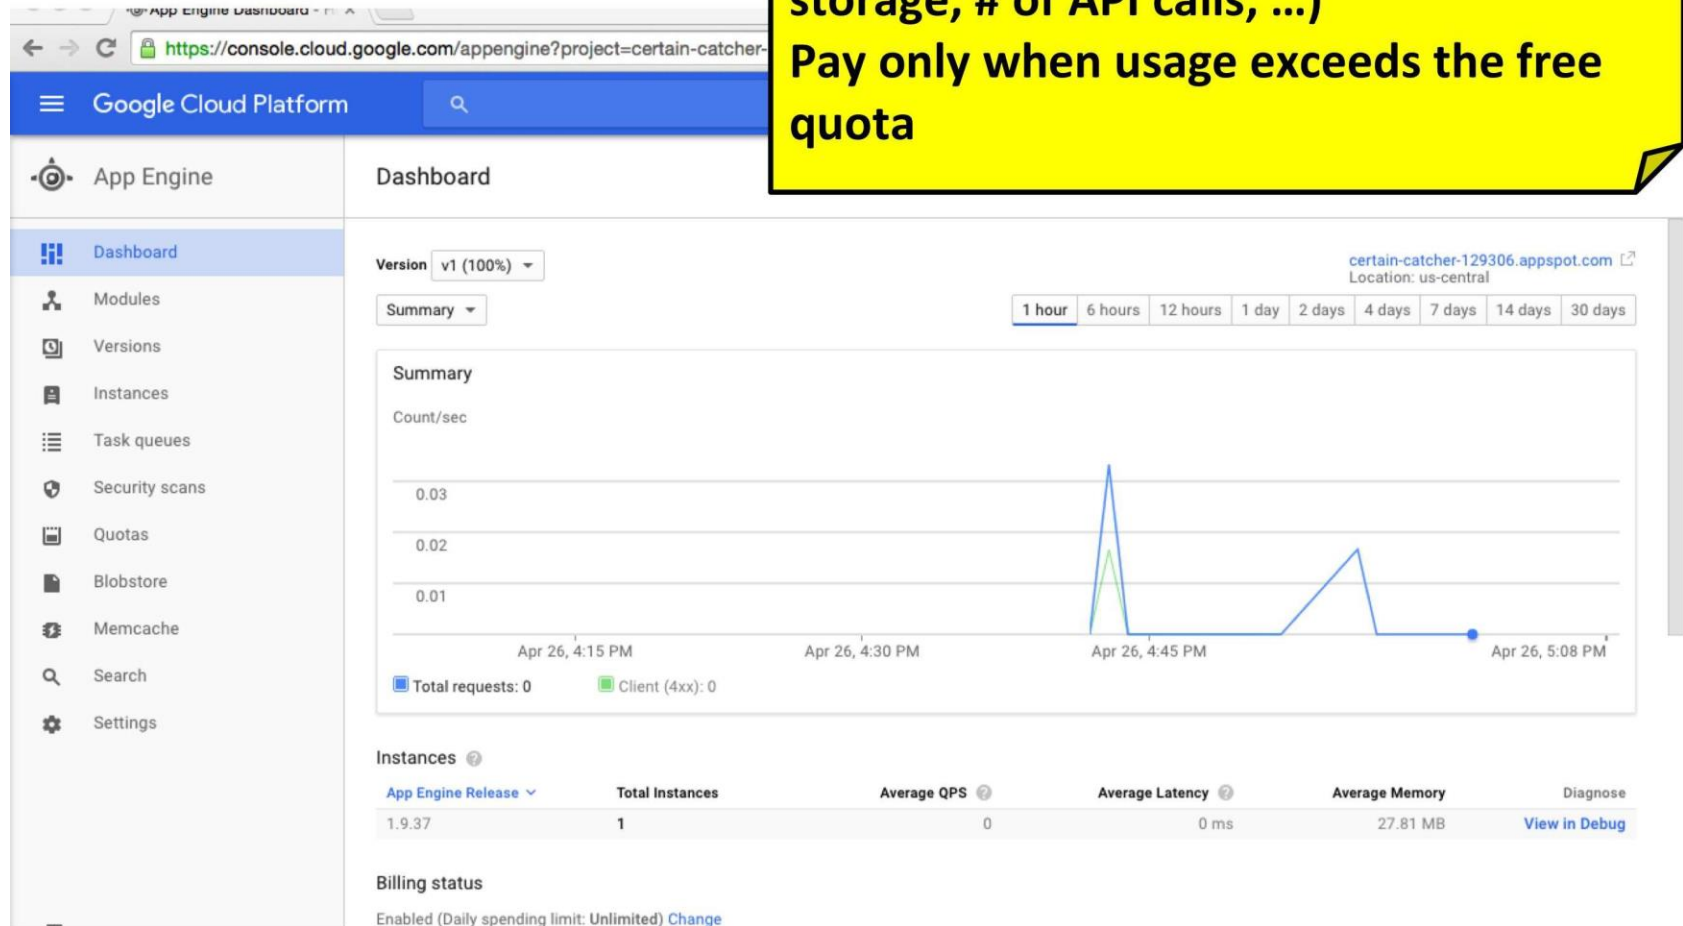
Scale up/down, load balancing,  
replication, database management, ...  
many services are provided by GAE.





# EXAMPLE 2: APP ENGINE

**6. Check your resource usage (CPU, storage, # of API calls, ...)  
Pay only when usage exceeds the free quota**





# SOFTWARE AS A SERVICE

## **Service provider provides:**

- Hardware infrastructure
- OS and platform software (middleware)
- Distributed storage management
- Load balancing, replication, migration
- Management and Monitoring services
- Application

## **Client provides:**

- Data



# Challenge

- Learn new application
- Deal with potential restrictions
- Web interface, restricted functionality
- No offline access, no local storage

## Challenges – Provider:

- Transparency (naming, redirection)
- Scalability: replication and load balancing decisions
- Synchronisation and coordination
- Security
- Fault tolerance
- Monitoring
- Software maintenance and sys admin
- Application development and maintenance

# KEY CHALLENGES OF CLOUD COMPUTING

## Scalability:

- Datacentre vs Global
- Partitioning
- Services and Data
- Replication


## Consistency:

- Dealing with consequences of CAP Theorem
- Dealing with un-usability of eventual consistency

# KEY CHALLENGES OF CLOUD COMPUTING

---

- SLA (Service Level Agreement): guarantees given by provider
  - How reliable are the guarantees?
  - What is the consequence if they aren't met?
- Redundancy and Replication
  - within same provider (e.g. Availability Zones, Regions, etc.)
  - migration across providers
- Geographically distributed architecture



# KEY CHALLENGES OF CLOUD COMPUTING

- ➔ **Design for failure: Chaos Monkey**
- **test how well system deals with failure**
- **regularly and randomly kill system services**



# Security and Privacy:

- → **External threats**
- **Denial of Service**
- **Infrastructure or platform service compromise**
- **SaaS compromise: data theft**
- → **Co-located threats: other customers**
- **Isolation: but, covert channels, bugs in isolation**
- → **Privacy: data collected by providers**
- **IaaS and PaaS providers: encryption only helps a bit**
- **SaaS providers: at mercy of service provider**
- **Governments and others: where is your data stored or processed?**

**Which laws apply?**

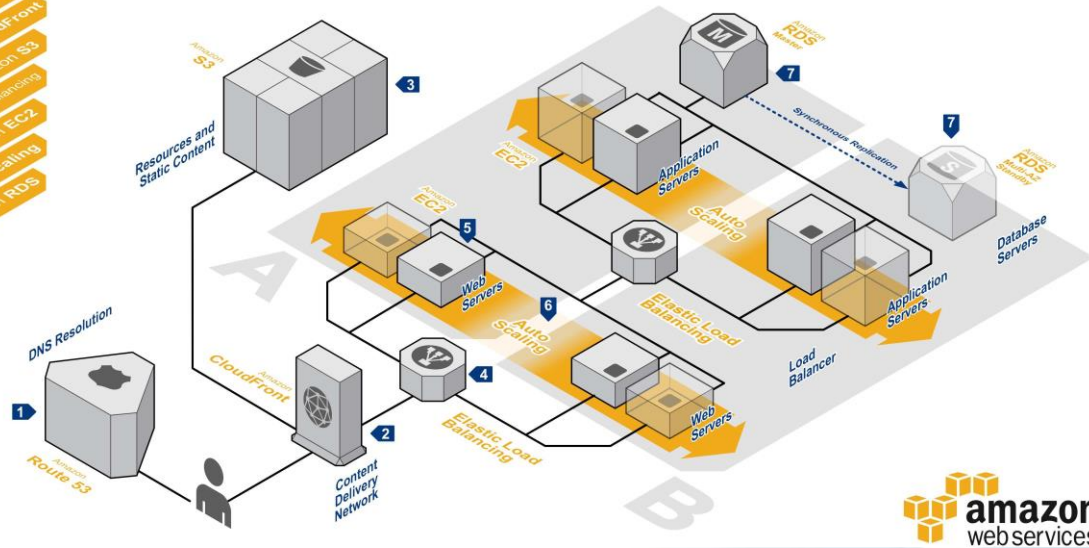
# DEVELOPING FOR THE CLOUD

Examples from Amazon:

**AWS Reference Architectures**  
Amazon Route 53  
Amazon CloudFront  
Amazon S3  
Amazon ElastiCache  
Amazon EC2  
Amazon Auto Scaling  
Amazon RDS

## WEB APPLICATION HOSTING

Highly available and scalable web hosting can be complex and expensive. Dense peak periods and wild swings in traffic patterns result in low utilization of expensive hardware. Amazon Web Services provides the reliable, scalable, secure, and high-performance infrastructure required for web applications while enabling an elastic, scale-out and scale-down infrastructure to match IT costs in real time as customer traffic fluctuates.



**amazon**  
web services

<http://aws.amazon.com/architecture/>